

Ministry of Higher Education
Faculty of Engineering
Computer and System department

Sort Algorithm – Parallel project

This project addresses the challenge of reducing execution time for computationally intensive problems by exploiting shared memory parallelism. A serial algorithm was transformed into a parallel version to improve performance while maintaining correctness in PARALLEL PROGRAMMING.

BY

Nada Salah Ali Ibrahim Ali (N/A)

SUPERVISORS

DR. Al-Mahdi

Alexandria. Egypt

December. 2025

How the Algorithm Works

This is a rank sort (or counting sort) algorithm. For each element $X[i]$, it determines its correct position in the sorted array by counting how many elements are smaller than it, with a tie-breaking rule for equal elements.

Step-by-step:

1. For each index i (0 to $N-1$):
2. Initialize $count = 0$
3. For each index j (0 to $N-1$):
4. Count how many elements $X[j]$ are strictly less than $X[i]$
5. For equal elements ($X[j] == X[i]$), count those with $j < i$ (earlier position)
6. Place $X[i]$ at position $count$ in the temporary array tmp
7. Copy tmp back to X

Pseudocode:

Serial: Loop executes sequentially.

```
void sequential_sort(std::vector& X) {  
  
    unsigned int i, j, count, N = X.size();  
    std::vector<unsigned int> tmp(N);  
  
    for (i = 0; i < N; i++) {  
        count = 0;  
        for (j = 0; j < N; j++) {  
            if (X[j] < X[i] || X[j] == X[i] && j < i)  
                count++;  
        }  
        tmp[count] = X[i];  
    }  
  
    std::copy(tmp.begin(), tmp.end(), X.begin()); }
```

Parallel: Loop parallelized using OpenMP pragma.

```

void sequential_sort(std::vector<int>& X) {

    unsigned int i, j, count, N = X.size();
    std::vector<unsigned int> tmp(N);
    # pragma omp parallel for private (i,j,count) shared(x,tmp,N)
    for (i = 0; i < N; i++) {
        count = 0;
        for (j = 0; j < N; j++) {
            if (X[j] < X[i] || X[j] == X[i] && j < i)
                count++;
        }
        tmp[count] = X[i];
    }

    std::copy(tmp.begin(), tmp.end(), X.begin()); }
```

Validation:

Normal code :

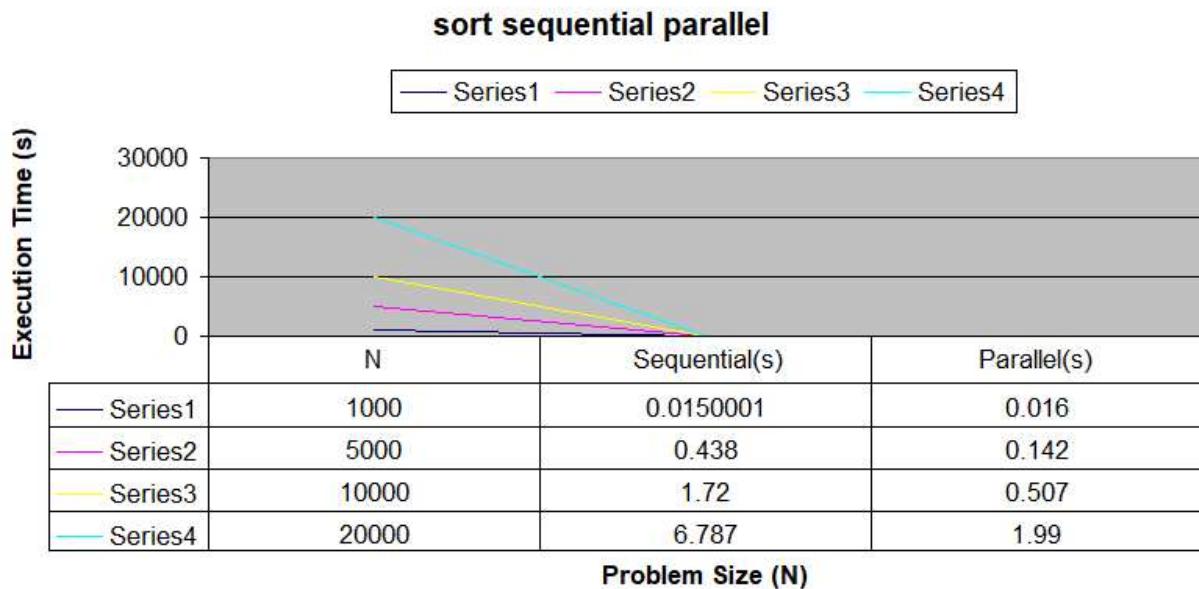
<https://github.com/NadaSalahAli/openmp-projects/blob/main/SeuentialSort.cpp>

<https://github.com/NadaSalahAli/openmp-projects/blob/main/parallelSort.cpp>

Different problem size for sequential and parallel

<https://github.com/NadaSalahAli/openmp-projects/blob/main/differentSizeSortSquParallel.cpp>

Graph :



Data Dependencies Analysis

Outer loop :

- No read-after-write and no write-after-write between iterations
- Outer loop iterations are independent and can run in parallel

Inner loop :

- There is a dependency (across j iterations)
- Inner loop has sequential reduction pattern
- If I had used parallel execution here, a race condition would have occurred

Shared Variables :

- tmp: Different iterations write to different positions (no conflict)
- count: Private to each outer iteration (should be thread-local)
- X: Read-only across threads (safe)
- N, i, j: Loop indices need proper handling

Performance Analysis:

Theoretical Complexity:

Sequential: $O(N^2)$ operations

Parallel (P threads): $O(N^2/P)$ operations and there is an overhead

Speedup & Efficiency:

Perfect speedup: $S(P) = P$ (ideal)

Actual speedup: Limited by:

- Memory bandwidth: All threads read the same array X repeatedly
- Cache contention: Multiple threads accessing same memory locations
- Load imbalance: Each iteration does the same work
- OpenMP overhead: Thread creation, scheduling, synchronization

Efficiency:

- For small N : Poor (overhead dominates)
- For large N : Approaches 1 (good scaling)
- Memory-bound for very large N : Efficiency drops

challenges :

1. To stay away from AI assistance
2. To choose a problem that I can work on without help
3. To try to raise the level of my human analysis of the algorithm

References:

[2] T. G. Mattson, B. A. Sanders, and B. L. Massingill, Parallel Programming: Concepts and Practice, Boston: Addison-Wesley, 2004, ch. 6, p. 218.